

Marc Berthelette
Smiths Falls, ON, Canada | marc.berthelette@gmail.com
linkedin.com/in/marc-berthelette-3aba6917

September 18, 2026

Elastic

Re: Senior Site Reliability Engineer (Observability & Analytics) — Platform Infra

Dear Elastic hiring team,

I am applying for the Senior Site Reliability Engineer position on Platform Observability & Analytics. The role describes the kind of reliability work I value most: operating infrastructure that tells engineers what production is actually doing, responding clearly when it fails, and turning each incident into a lasting improvement rather than another page. I bring 18 years across systems, software, DevOps, and platform engineering, including production ELK, Linux fleets, Terraform, Ansible, Python, Kubernetes, GitOps, and 24/7 on-call.

At Tink, I supported more than 600 Linux servers and participated in the on-call rotation for all production systems and applications. Through root-cause analysis and prevention, paging frequency fell from roughly one page every two rotations to one every four or five. One recurring failure involved an unstable Elasticsearch cluster; I connected monitoring to Bash remediation so that known condition restored itself instead of waking an engineer. I also implemented centralized Elasticsearch, Logstash, and Kibana logging and managed fleet configuration declaratively with Puppet and Spacewalk.

At Desjardins, I designed and built a Python/FastAPI service that tested the health and configuration of approximately 250 enterprise Artifactory repositories every four hours and published the results to Dynatrace. I made the validation logic tool-agnostic and the test set configurable, allowing the same system to become the regression safety net for an enterprise platform upgrade. I also automated isolated environments with Kubernetes, ArgoCD, and Helm and built Splunk alerts for changes made outside approved delivery paths.

At Énergir, I now drive delivery and reliability improvements across 12 teams and more than 70 applications. I lead root-cause analysis for critical incidents, build monitoring and automated remediation, and use Terraform and Ansible to make environments repeatable. The work also requires the behaviours Elastic calls out explicitly: owning ambiguous problems end to end, reviewing production changes carefully, documenting operating practices, raising risks, and enabling engineers outside my reporting line.

I want to be precise about the remaining gaps. I have not operated Elastic Cloud's ECH, ECE, or ECK orchestration layers, and Go is not part of my demonstrated production experience. My Kubernetes experience is practical—automated environments, ArgoCD, Helm, and platform validation—rather than deep cluster-engineering ownership. I would bring those boundaries honestly, along with unusually direct experience in the rest of the role: Elasticsearch operations, production observability, Linux, Python, Terraform, configuration management at fleet scale, on-call, RCA, regulated infrastructure, and automation that removes repeat pages.

I would welcome the opportunity to help operate and harden the observability infrastructure that Elastic Cloud engineers rely on across more than 200 hosted deployments.

Sincerely,

Marc Berthelette