

Dear Nexxa.ai Hiring Team,

I'm applying for the Staff DevOps Engineer role. Owning infrastructure end-to-end, across cloud, on-prem, and edge, for a small and fast-moving team is the kind of scope I've consistently taken on throughout my career, even when it meant learning unfamiliar territory under a delivery deadline.

At Énergir, I led cloud modernization work migrating several mission-critical services onto Microsoft Azure, building Terraform and Ansible infrastructure as code that eliminated configuration drift across environments, and drove the root-cause analysis and automated remediation that reduced MTTR across the application estate. At Desjardins, I automated on-demand infrastructure provisioning using Kubernetes and ArgoCD and built a Python-based service that continuously validated infrastructure configuration at scale, extending the engagement past its original scope by making the underlying logic reusable rather than one-off. Earlier at Tink, I conceived and built, unprompted, a fleet-wide software inventory platform surfacing security and obsolescence risk to product owners who had no prior visibility into it — the kind of independent project leadership this role asks for.

I want to be direct about two specific gaps: I haven't worked with GPU workload scheduling or ML/AI infrastructure (model training and serving pipelines, feature stores), and my production language experience is Python, Bash, C++, and Java rather than Go. What I bring instead is a demonstrated pattern of learning unfamiliar platforms cold and fast — Salesforce, Kubernetes, and a full AWS build environment, each with no prior experience going in — and I'd apply that same approach to closing the GPU/ML-infrastructure gap quickly.

I'd welcome the opportunity to discuss how this background applies to Nexxa's infrastructure needs. Thank you for your consideration.

Sincerely, Marc Berthelette